

Introduction(week 1, part 3)

11 Jan 2023

Table of contents

books	1
Fisher's irises	2
linear models	2
least squares	2
regression: nuts and bolts	3
regression as classification	3
nearest-neighbor	4
consistency	4
from nearest to k -NN	4
from NN to kernel smoothers (ADA, ESL 2.8.2)	4
dimensionality	5
bias-variance expansion/trade-off	5
the picture	6
effective degrees of freedom	6
references	6

books

- ESL and ADA cover very similar material
- both compare linear regression and nearest-neighbour methods as opposite ends of a complexity spectrum
- ESL has more on dimensionality

Fisher's irises

- Canadian content (irises of the Gaspé peninsula) (Fisher 1936)
- Fisher was a eugenisit (!) (Bodmer et al. 2021)
- multiple versions of the data set, with errors ... (Bezdek et al. 1999)
- alternative: [Palmer penguins dataset](#)

Fisher, R. A. 1936. "The Use of Multiple Measurements in Taxonomic Problems." *Annals of Eugenics* 7 (2): 179–88. <https://doi.org/10.1111/j.1469-1809.1936.tb02137.x>.

Bodmer, Walter, R. A. Bailey, Brian Charlesworth, Adam Eyre-Walker, Vernon Farewell, Andrew Mead, and Stephen Senn. 2021. "The Outstanding Scientist, R.A. Fisher: His Views on Eugenics and Race." *Heredity* 126 (4): 565–76. <https://doi.org/10.1038/s41437-020-00394-6>.

Bezdek, J. C., J. M. Keller, R. Krishnapuram, L. I. Kuncheva, and N. R. Pal. 1999. "Will the Real Iris Data Please Stand Up?" *IEEE Transactions on Fuzzy Systems* 7 (3): 368–69. <https://doi.org/10.1109/91.771092>.

linear models

- can write out as $\hat{Y} = \hat{\beta}_0 + \sum X_j \hat{\beta}_j$
- go almost immediately to $\hat{Y} = X^\top \hat{\beta}$ or $\langle X, \beta \rangle$ or $\mathbf{X}\beta$
- \mathbf{X} is the *model matrix* (sometimes "design matrix")
- usually includes an intercept column
- can contain *any* (precomputed) functions of input variables
- input vars (directly measured) \rightarrow predictor vars (transformations, basis expansions, etc.)
- 1D examples

least squares

- choose L2 norm (p-norm = $(\sum |x|^p)^{1/p}$)
- $\sum_i (Y_i - X_i \beta)^2$
- equivalent to $(\mathbf{y} - \mathbf{X}\beta)^\top (\mathbf{y} - \mathbf{X}\beta)$
- differentiate and solve: $\hat{\beta} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$
- **hat matrix:**

$$\begin{aligned}\hat{x} &= H\mathbf{y} = \mathbf{X}\hat{\beta} \\ &= \mathbf{X}(\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y} \\ H &= \mathbf{X}(\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top\end{aligned}$$

- regression as a *linear filter*
- cf. explicit expression in ADA 1.52

regression: nuts and bolts

- never do naive linear algebra!
 - `fortunes::fortune("SL0000W")`
- R: QR decomp with Householder rotations
 - see: [lm.c](#), [dqlrs](#), [dqrdc2](#) (from the beginning)
- Julia ???
- will dig into computational details a bit next week (?)

regression as classification



Figure 1: fig2.1

- slightly weird
- looseness of “classification” vs “regression”
 - should probably use *discriminant analysis* here
 - or logistic regression

nearest-neighbor

- $\frac{1}{k} \sum_{x_i \in N_k(x)} y_i$
- also a linear smoother: columns of \mathbf{X} are $1/k \times$ *indicator variables*: ($x_i \in N_k(x)$) and the β values are y_i
- **hardening** predictions: better to leave as a probability?

consistency

- if we want *consistency* we need (roughly) the number of observations used to make a prediction to grow (fast enough) with the total sample size N
- doesn't hold for fixed k
- but will work if $k/n \rightarrow 0$ as $k \rightarrow \infty, n \rightarrow \infty$ (ADA)

from nearest to k -NN

- unlike (this version of) linear regression, complexity is adjustable
- from “nearest neighbor” to “ N -n.n.” (i.e. the mean)
- tuning parameter (k) is **discrete** (e.g. awkward for optimization)

poll: what are some ways we can modify linear regression to have adjustable complexity?

from NN to kernel smoothers (ADA, ESL 2.8.2)

- kernel **density estimation** may be familiar
- generalize “nearest neighbor” kernel
- from ADA:

$$\hat{\mu}(x) = \sum_i \left(\frac{K(x_i, x)}{\sum_j K(x_j, x)} \right) y_i$$

- or $K(d(x_i, x)/h)$ where h is the **bandwidth**
- also not consistent unless we let $h \rightarrow 0$ and $n \rightarrow \infty$ (at the “right” rate)

dimensionality

- **curse of dimensionality**
- more points are near the edge of a set
- more points are needed to “fill in”/characterize a density
- the mode of a distribution is no longer “typical” in some sense (Gelman 2020)
- heuristic: surface to volume ratio of a p -ball is p^1
- results from ESL: distance from origin to nearest point (of N) in dimension p

$$d(p, N) = \left(1 - \frac{1}{2}\right)^{1/p}$$

- Bayesians: the mode is not “typical” Gelman (2020)

bias-variance expansion/trade-off

- general expansion of $E[(y - \hat{\mu}(x))^2]$, expanded as

$$E\left[\underbrace{(y - \mu(x))}_{\text{diff betw } y, \text{ true RF}} + \underbrace{(\mu(x) - \hat{\mu}(x))}_{\text{diff betw true RF, chosen approx}}\right]^2$$

- take expectations, drop 0 terms \rightarrow variance + bias²
- allow for variation across training sets, get $\sigma^2 + \text{var} + \text{bias}^2$ (ESL 2.47)
- ADA example: true function (sine) is *worse than a constant function* for noisy data (Fig 1.3). cf. Walters and Ludwig (1981)
- bias \approx **within-sample** error; easy to minimize

Gelman, Andrew. 2020. “The Typical Set and Its Relevance to Bayesian Computation.” *Statistical Modeling, Causal Inference, and Social Science*. <https://statmodeling.stat.columbia.edu/2020/08/02/the-typical-set-and-its-relevance-to-bayesian-computation/>.

¹ surprisingly slow!

Gelman, Andrew. 2020. “The Typical Set and Its Relevance to Bayesian Computation.” *Statistical Modeling, Causal Inference, and Social Science*. <https://statmodeling.stat.columbia.edu/2020/08/02/the-typical-set-and-its-relevance-to-bayesian-computation/>.

Walters, Carl J., and Donald Ludwig. 1981. “Effects of Measurement Errors on the Assessment of Stock–Recruitment Relationships.” *Canadian Journal of Fisheries and Aquatic Sciences* 38 (6): 704–10. <https://doi.org/10.1139/f81-093>.

the picture

2. Overview of Supervised Learning




FIGURE 2.11. Test and training error as a function of model complexity.

- want to find the ‘sweet spot’ with low computational effort, with minimal assumptions, and without snooping

effective degrees of freedom

- may be able to compute a complexity measure, for simple cases
- ... such as linear weights
- **trace of the hat matrix**; ADA 1.66-1.68
- related to ESL 2.28

references