

# Non-Gaussian responses (week 3?)

31 Jan 2023

## Table of contents

Non-Gaussian responses . . . . .	1
Some answers . . . . .	2
Why not linear? . . . . .	2
Logistic regression (ESL § 4.4) . . . . .	2
Log-likelihood . . . . .	3
Newton step . . . . .	3
Newton vs IRLS . . . . .	3
Families . . . . .	4
Regularized versions . . . . .	4
revisiting ridge by data augmentation . . . . .	4
ridge + IRLS . . . . .	5
proximal gradient descent/Newton . . . . .	5
proximal operator . . . . .	5
proximal Newton (IRLS) . . . . .	6
more computational details . . . . .	6
sparse model matrices (side note) . . . . .	7

## Non-Gaussian responses

- Why worry about it?
- Isn't least-squares good enough?
- **poll** ([polleverywhere](#))

## Some answers

- heteroscedasticity (Gauss-Markov only applies to homog. variance)
- still unbiased but no longer minimum variance
- maybe we shouldn't (e.g. **linear probability model** in econometrics)
  - adjust for heteroscedasticity with **robust/sandwich estimators** etc. (White):

$$\hat{\mathbf{V}} = (\mathbf{X}^\top \mathbf{X})^{-1} (\mathbf{X}^\top \mathbf{G} \mathbf{X}) (\mathbf{X}^\top \mathbf{X})^{-1}$$

where  $\mathbf{G} = \text{Diag}(\hat{\varepsilon}_i^2)$  (contrast with  $s^2(\mathbf{X}^\top \mathbf{X})^{-1}$ )

- if we have
- if we have **nonlinear** models, MLEs are no longer unbiased

## Why not linear?

- actual nonlinear patterns (but can handle these by transformation/basis expansion)
- unrealistic predictions (e.g. probabilities outside of  $[0, 1]$ )
- varying effects (e.g. effect of a 1-unit change in  $x$  on probability must differ depending on baseline probability)
- Why not transform? **poll** ([polleverywhere](#))

## Logistic regression (ESL § 4.4)

- Worst-case scenario (farthest from Gaussian)
- ESL starts with a *multinomial* model:

$$\log \left( \frac{\Pr(G = i | X = x)}{\Pr(G = K | X = x)} \right) = \beta_{i0} + \beta_i^\top x, \quad i \in 1 \dots K-1$$

(and so  $\Pr(G = K | X = x) = 1 / \left( 1 + \sum_{i=1}^{K-1} \exp(\beta_{i0} + \beta_i^\top x) \right)$ )

- independent of baseline/reparameterization
- log-likelihood  $\sum \log p_i(x_i; \theta)$  where  $\theta$  is the complete set of parameters

## Log-likelihood

- for two categories, log-likelihood simplifies to

$$\begin{aligned} & \sum (y_i \beta^\top x_i - \log(1 + e^{\beta^\top x_i})) \\ &= \sum (y_i \eta_i - \log(1 + \exp(\eta_i))) \end{aligned}$$

- **weight matrix**  $\mathbf{W} = \text{Diag}(p(1-p))$ 
  - more generally,  $\text{Diag}(1/\text{Var}(\mu))$
- score equation:
  - $\sum_i = 1^N x_i (y_i - p(x_i; \beta))$
  - Newton update is  $\beta^* - \mathbf{H}^{-1} \mathbf{g}$
  - gradient:  $\mathbf{X}^\top (\mathbf{y} - \mathbf{p})$
  - generally  $\mathbf{X}^\top (\mathbf{y} - \mu) = \mathbf{X}^\top (\mathbf{y} - g^{-1}(\eta))$
- Hessian:  $-\mathbf{X}^\top \mathbf{W} \mathbf{X}$
- solution is  $(\mathbf{X}^\top \mathbf{W} \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{W} \mathbf{z}$
- where  $\mathbf{z} = \mathbf{X} \beta_0 + \mathbf{W}^{-1}(\mathbf{y} - \mathbf{p})$  is the adjusted response

## Newton step

- **iteratively reweighted least squares**
- solve

$$\mathbf{X}^\top \mathbf{W} \mathbf{X} \beta^* = \mathbf{X}^\top \mathbf{W} \mathbf{z}$$

- [4C03 notes](#)
- [4C03 notes 2](#)

## Newton vs IRLS

- Newton vs *Fisher scoring* (expected value of the Hessian); equivalent for the *canonical link* (e.g. logistic for binary data, log for Poisson data)
- link mostly important for interpretation
- can be disregarded (?) if we are going to handle nonlinearity by basis expansion
- convergence? (Mount 2012)

Mount, John. 2012. “How Robust Is Logistic Regression?” *Win Vector LLC*. <https://win-vector.com/2012/08/23/how-robust-is-logistic-regression/>.

## Families

- Gaussian, Poisson, binomial (binary)
- May need to compute *scale/dispersion* parameter
  - for exponential families, calculate as  $\sqrt{D/(n-p)}$  where  $D$  is the *deviance* (-2 log likelihood, equal to SSQ for Gaussian)
  - not exactly the MLE but good enough
- **over-dispersion**: quasi-likelihood
- more complex families (negative binomial etc.) have an additional, non-collapsible parameters, need to estimate by MLE (or **profiling**)

## Regularized versions

- lasso, ridge, or elasticnet
- score equations:  $\mathbf{x}^{\text{top}}(\mathbf{y} - \mathbf{p}) = \lambda \cdot \text{sign}(\beta_j)$  for **active** variables (non-zero coeffs)

## revisiting ridge by data augmentation

- we want to minimize  $\|\mathbf{y} - \mathbf{X}\beta\|_2^2 + \lambda\|\beta\|_2^2$
- the solution to the original regression equations was  $\hat{\beta} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$
- Set

$$\mathbf{B} = \begin{pmatrix} \mathbf{X} \\ \sqrt{\lambda} \mathbf{I} \end{pmatrix}$$

- ridge regression should still be solvable by data expansion, i.e. in the IRLS loop use

$$\mathbf{B} = \begin{pmatrix} \mathbf{X}^\top \mathbf{W} \mathbf{X} \\ \sqrt{\lambda} \mathbf{I} \end{pmatrix}$$

and

$$\mathbf{y}^* = (\mathbf{y} \ 0)$$

- so that  $\mathbf{B}^\top \mathbf{B} = \mathbf{X}^\top \mathbf{X} + \lambda \mathbf{I}$  and the  $\mathbf{X}^\top \mathbf{y}$  term is unchanged

## ridge + IRLS

- recall that we need to iteratively solve

$$\mathbf{X}^\top \mathbf{W} \mathbf{X} \beta^* = \mathbf{X}^\top \mathbf{W} \mathbf{z}$$

- if we want to solve the **weighted** least-squares problem from IRLS, we would normally take the QR decomposition of  $\mathbf{X}' = \mathbf{X} \sqrt{\mathbf{W}}$  (so that  $\mathbf{X}'^\top \mathbf{X}' = \mathbf{X}^\top \mathbf{W} \mathbf{X}$ )
- enhance this by adding  $\sqrt{\lambda} I$  to  $X$  and zeros to  $\mathbf{z}$  (no longer  $\mathbf{y}$ ) as before
- ¿ try out [enhanced GLM](#) ?

## proximal gradient descent/Newton

- solving the optimization problem for non-differentiable penalties
- previous solution (cyclic coordinate descent)
- simpler strategies (cyclic coordinate descent) may not work as well
- **proximal** gradient descent or **proximal** IRLS:
- like the pathwise coordinate descent solution from lasso:
- **solution:**

$$\tilde{\beta}_j(\lambda) \leftarrow S \left( \sum_{i=1}^N x_{ij} (y_i - \tilde{y}_i^{(j)}), \lambda \right)$$

- where  $S(t, \lambda) = \text{sign}(t)(|t| - \lambda)$
- except that we can no longer jump straight to the correct solution.

## proximal operator

- separate objective function into **smooth** part (likelihood/RSS/etc., **plus** ridge penalty) and **non-smooth** part (typically an L1 regularization term)
- **proximal operator:**

$$\text{argmin}_u \left( \underbrace{h(u)}_{\text{nonconvex part}} + \frac{1}{2} \|u - x\|_2^2 \right)$$

- for  $h = \lambda \|\beta\|_1$  (lasso penalty), we get the soft-threshold operator

$$\begin{cases} \beta_i - \lambda & \text{if } \beta_i < -\lambda \\ 0 & \text{if } -\lambda < \beta_i < \lambda \\ \beta_i + \lambda & \text{if } \beta_i > \lambda \end{cases}$$

(from [Ryan Tibshirani's notes on optimization](#))

## proximal Newton (IRLS)

- take a Newton/IRLS step
- apply the prox operator to soft-threshold
- not going to get into the details! thresholding is more complex than the gradient descent rule (Lee, Sun, and Saunders 2014)
- need to solve

$$\operatorname{argmin}_u \left( \underbrace{h(u)}_{\text{nonconvex part}} + \frac{1}{2}(u - x)^\top H(u - x) \right)$$

i.e. replace  $\|u - x\|_2^2$  with a corresponding quadratic form

- $\zeta$  not sure if the solution corresponds easily to soft-thresholding again?
- also need to be careful about backtracking if necessary
- i.e. taking a Newton step  $-\mathbf{lg}$  is better than an uninformed gradient step  $t\mathbf{g}$  (where  $t$  is the **learning rate**) but might overshoot
- i.e., **don't try this at home**

([Ryan Tibshirani again](#))

Lee, Jason D., Yuekai Sun, and Michael A. Saunders. 2014. "Proximal Newton-Type Methods for Minimizing Composite Functions." *SIAM Journal on Optimization* 24 (3): 1420–43. <https://doi.org/10.1137/130921428>.

## more computational details

- from [glmnet family docs](#)
- using the *name* of the family ("poisson" etc.) uses hard-coded internal algorithms
  - faster (but scaling isn't too bad??)
  - less flexible (alternative families [links, variance functions])
  - slightly less robust (doesn't do backtracking)

### **sparse model matrices (side note)**

- expanding factors (categorical variables) may make  $p$  gigantic
- each factor  $f_i$  with  $n_i$  levels will be expanded via treatment contrasts, so  $\mathbf{X}$  will have (at least)  $\sum_i (n_i - 1)$  columns
- `glmnet::makeX, Matrix::sparse.model.matrix()`